



Let's talk MapReduce

Agenda

Before Map Reduce...

Map Reduce Overview

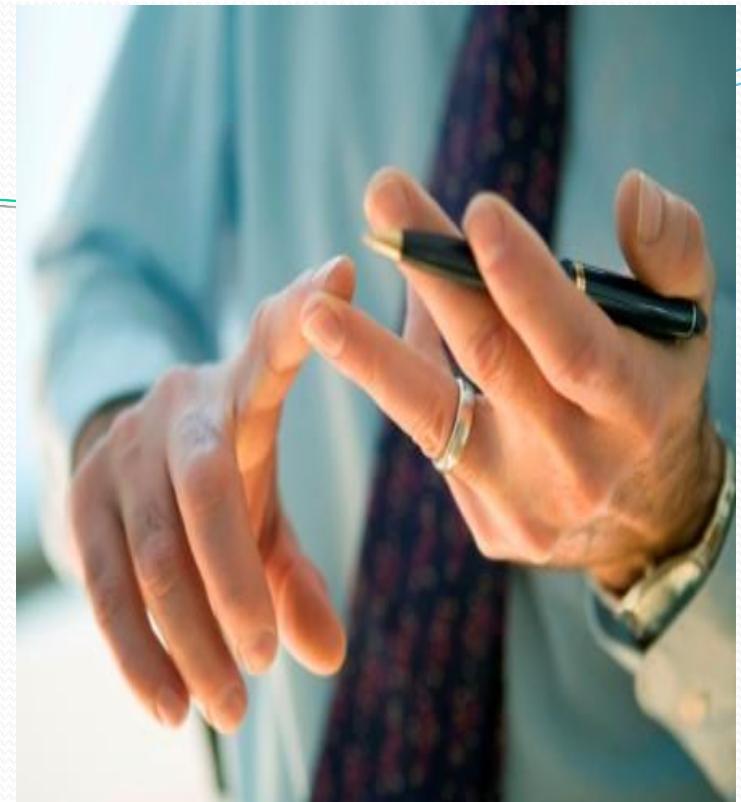
Word Count Problem

Word Count Flow and Solution

Map Reduce Flow

Algorithms for few Simple Problems

Algorithms for few Complex Problems



Before MapReduce...

- Large scale data processing was difficult!
 - Managing hundreds or thousands of processors
 - Managing parallelization and distribution
 - I/O Scheduling
 - Status and monitoring
 - Fault/crash tolerance
- MapReduce provides all of these, easily!

Overview of MapReduce

- Hadoop Map-Reduce is a software framework for easily writing applications which process vast amounts of data (multi-terabyte data-sets) in-parallel on large clusters (thousands of nodes) of commodity hardware in a reliable, fault-tolerant manner.
 - Programming model used by Google
 - A combination of the **Map** and **Reduce** models with an associated implementation
 - MapReduce is highly scalable and can be used across many computers and answers all problems mentioned earlier.
 - Many small machines can be used to process jobs that normally could not be processed by a large machine.
 - Programming Languages Supported: Java
 - Hadoop Streaming: Python, Ruby, Perl etc
 - Hadoop Pipes: C++

MapReduce Used for...

- At Google:
 - Index building for Google Search
 - Article clustering for Google News
 - Statistical machine translation
- At Yahoo!:
 - Index building for Yahoo! Search
 - Spam detection for Yahoo! Mail
- At Facebook:
 - Data mining
 - Ad optimization
 - Spam detection

Mapper

- Mapper reads the data in key/value pairs
- Outputs zero or more key/value pairs

map(in_key, in_value) -> (inter_key, inter_value) list

- Mapper may use or ignore key
 - The key is the byte offset into the file at which the line starts
 - The value is the contents of the line itself
 - Typically the key is considered irrelevant
- If the Mapper writes anything out, the output must be in the form of key/value pairs

Example: Turn into upper case

let map(k, v) = emit(k.toUpperCase(), v.toUpperCase())

Input - ('hadoop', 'big data') -> output - ('HADOOP', 'BIG DATA')

Reducer

- After Map phase, all intermediate values are combined together for given intermediate key into a *list*.
- All values associated with a particular intermediate key are guaranteed to go to the same Reducer
- The intermediate keys, and their value lists, are passed to the Reducer in sorted key order
- This step is known as the 'shuffle and sort'
- Reducer outputs zero or more key/value pairs

• Example:

```
let reduce(k, vals)
    sum = 0
    foreach int i in vals:
        sum += i
    emit(k, sum)
Input - ('bee', [111, 100, 110]) ->
Output -('bee', 321)
```

Identity Reducer:
let reduce(k, vals) =
 foreach v in vals:
 emit(k, v)

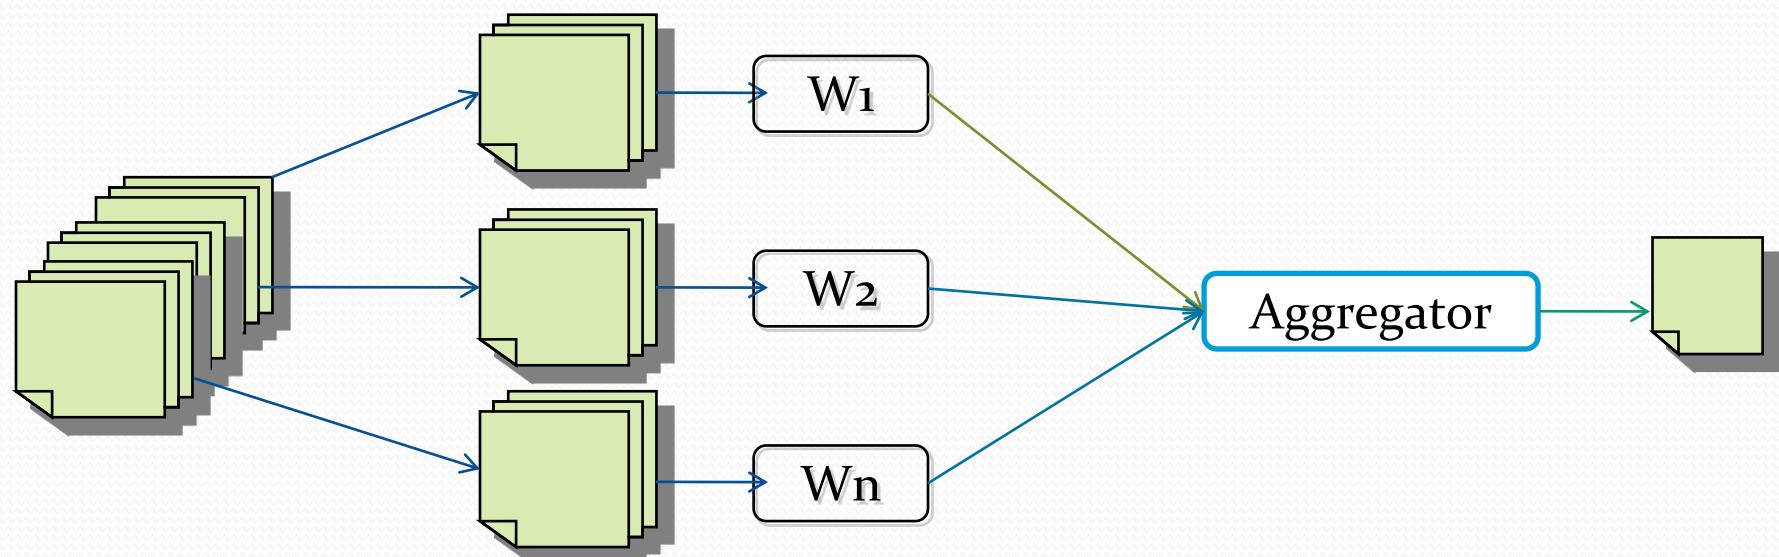
Input - ('abc', [111, 222, 333]) -> Output
('abc', 111), ('abc', 222), ('abc', 333),

How many Maps/Reducers?

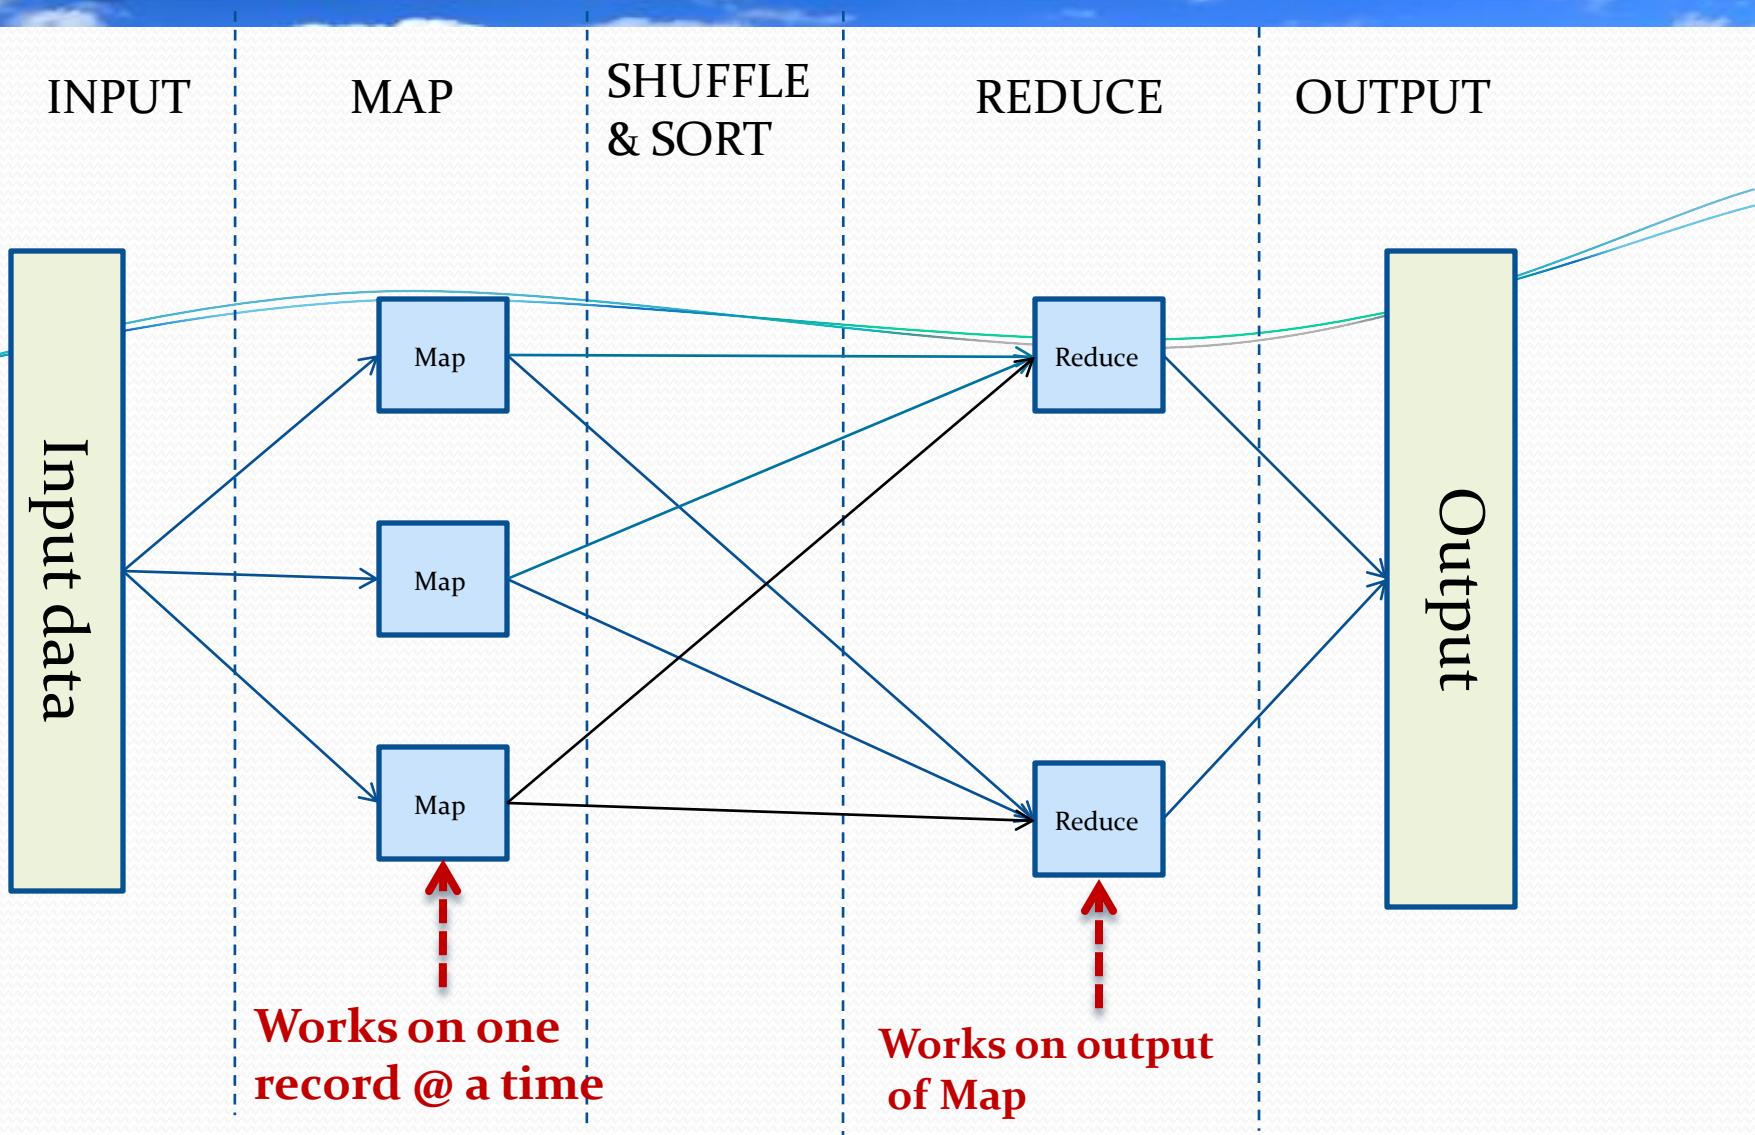
- Maps
 - Usually as many as the number of HDFS blocks being processed, this is the default
 - Else the number of maps can be specified as a hint
 - The number of maps can also be controlled by specifying the *minimum split size*
 - The actual sizes of the map inputs are computed by:
 - $\max(\min(block_size, data/\#maps), min_split_size)$
- Reduces
 - Unless the amount of data being processed is small
 - $0.95 * num_nodes * mapred.tasktracker.tasks.maximum$

Word Count - Distributed Solution

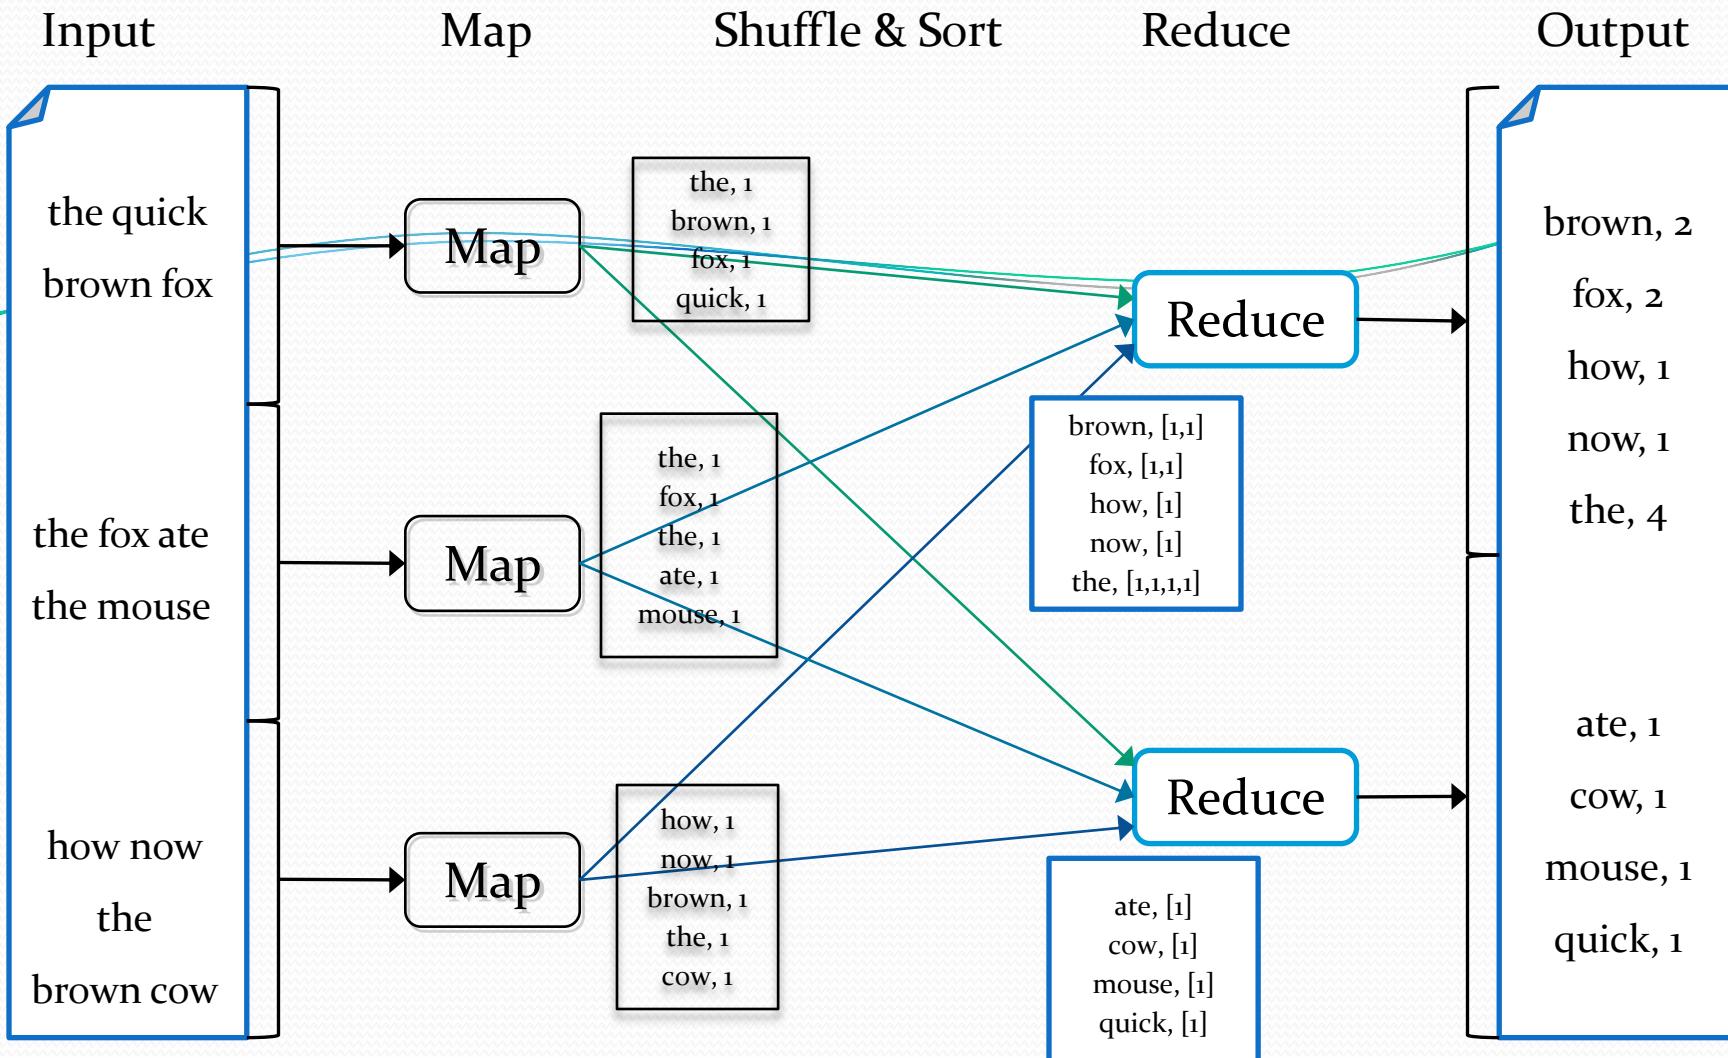
- Problem:
 - Large quantity of documents
 - Count the number of times each distinct word occurs in the documents
- Solution:
 - Divide and Conquer



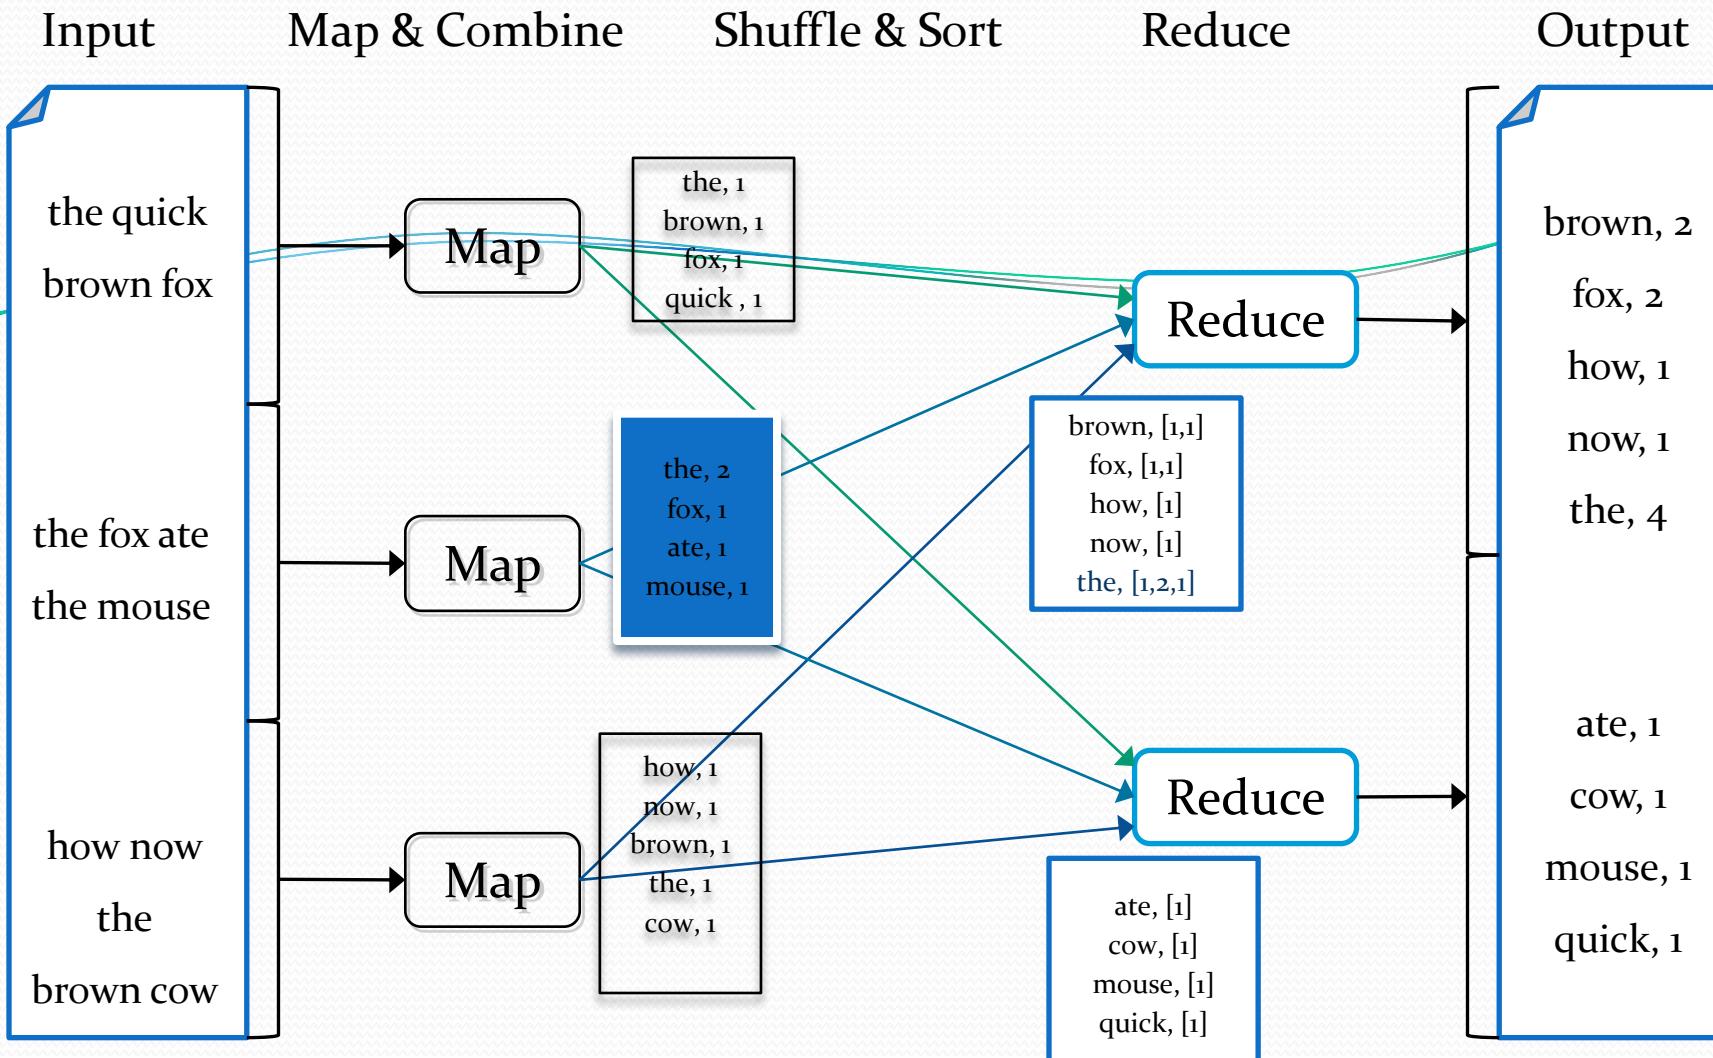
Map Reduce Flow



Word Count - Distributed Solution



Word Count - Distributed Solution



Word Count - Distributed Solution

```
def mapper(line):  
    foreach word in line.split():  
        output(word, 1)
```

```
def reducer(key, values[]):  
    output(key, sum(values))
```

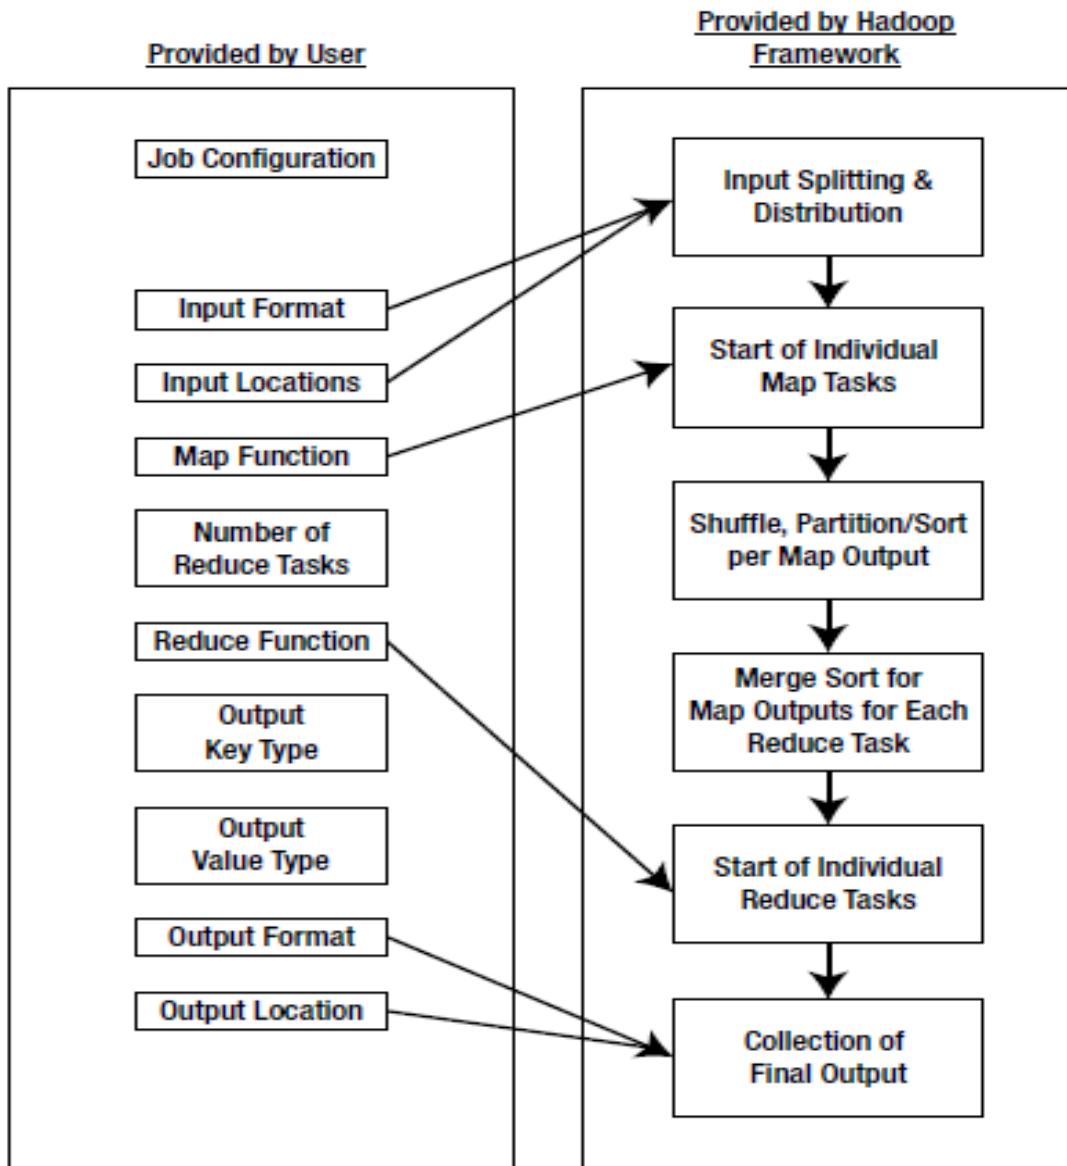
Map: $(k_1, v_1) \rightarrow \text{list } (k_2, v_2)$
Reduce: $(k_2, \text{list } v_2) \rightarrow \text{list } (k_3, v_3)$

Word Count - Thoughts

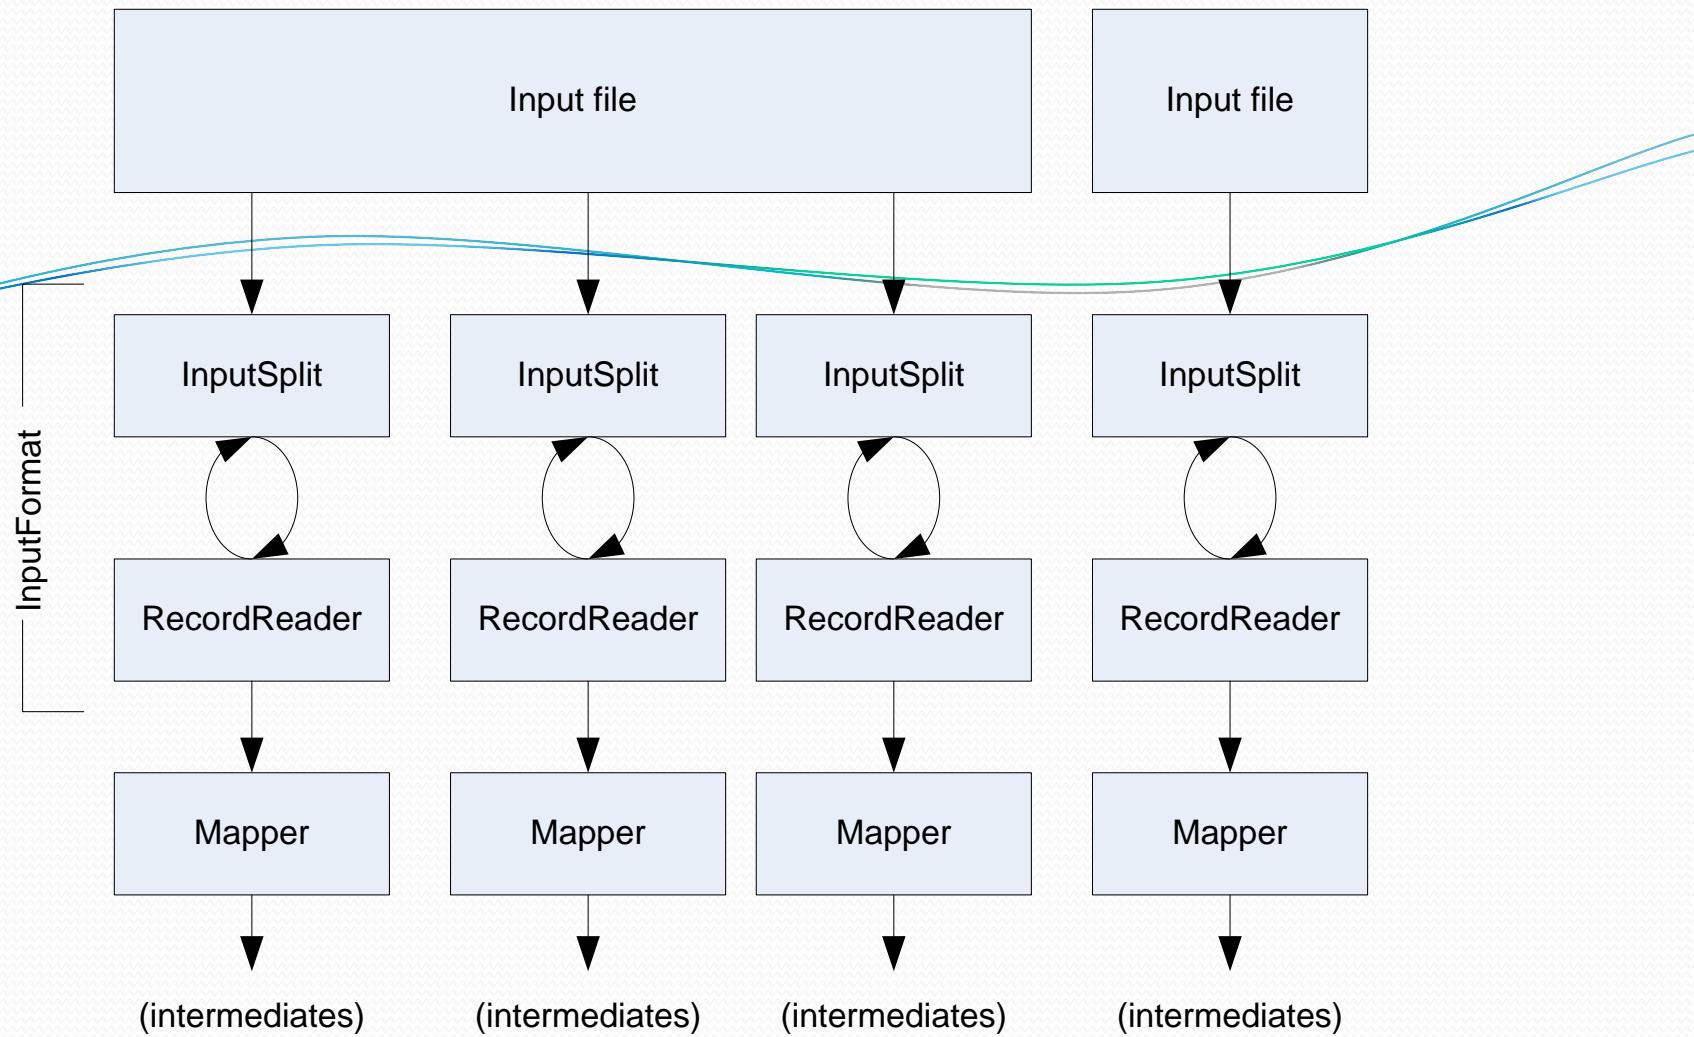
- Resource Requirements:
 - Use of commodity machines than super computers
 - Lot of storage space
 - Compute near data
 - Need a reliable coordinator
 - Distribute the Data
 - Coordinate the tasks
- Challenges:
 - A task on machine can crash, A task on machine can run very slow, A machine can crash
 - There may be data loss while data being moved
 - Coordinator has to be very reliable
 - Corpus division amongst nodes needs to be fairly equal

**WORD COUNT WAS A SIMPLE PROBLEM! IF
SOMEBODY TAKES CARE OF THESE
STANDARD THINGS, THEN IT CAN BE STILL
KEPT SIMPLE**

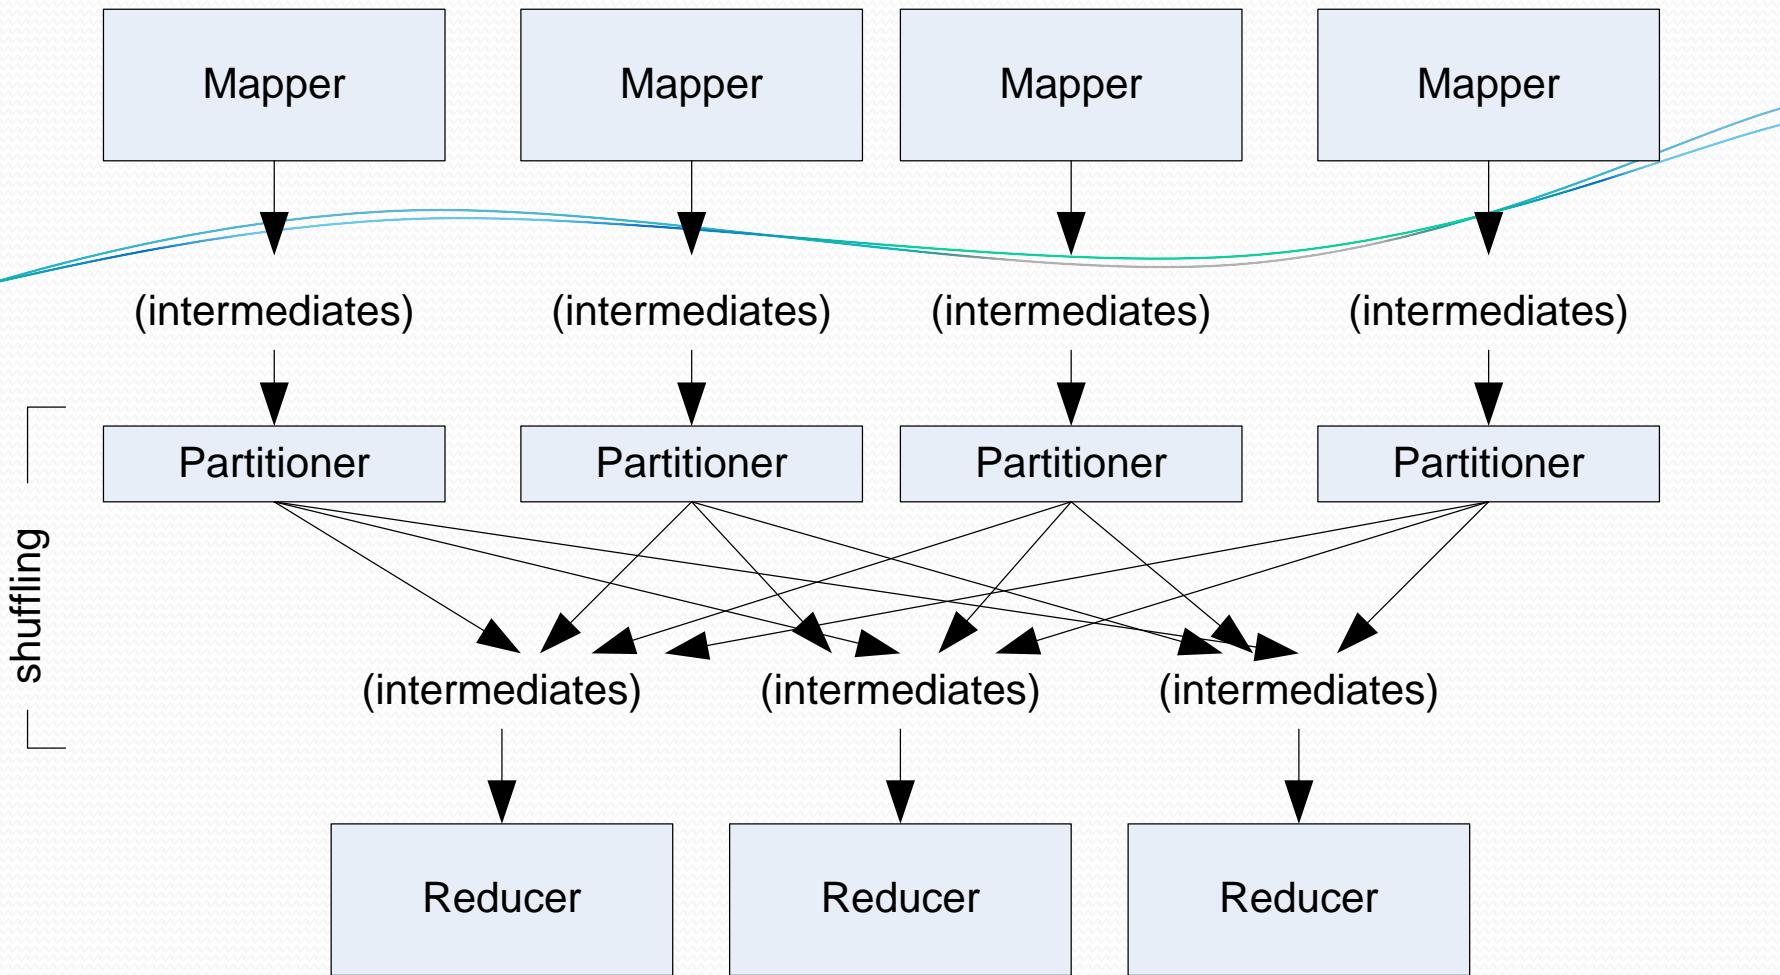
Parts of MapReduce Job



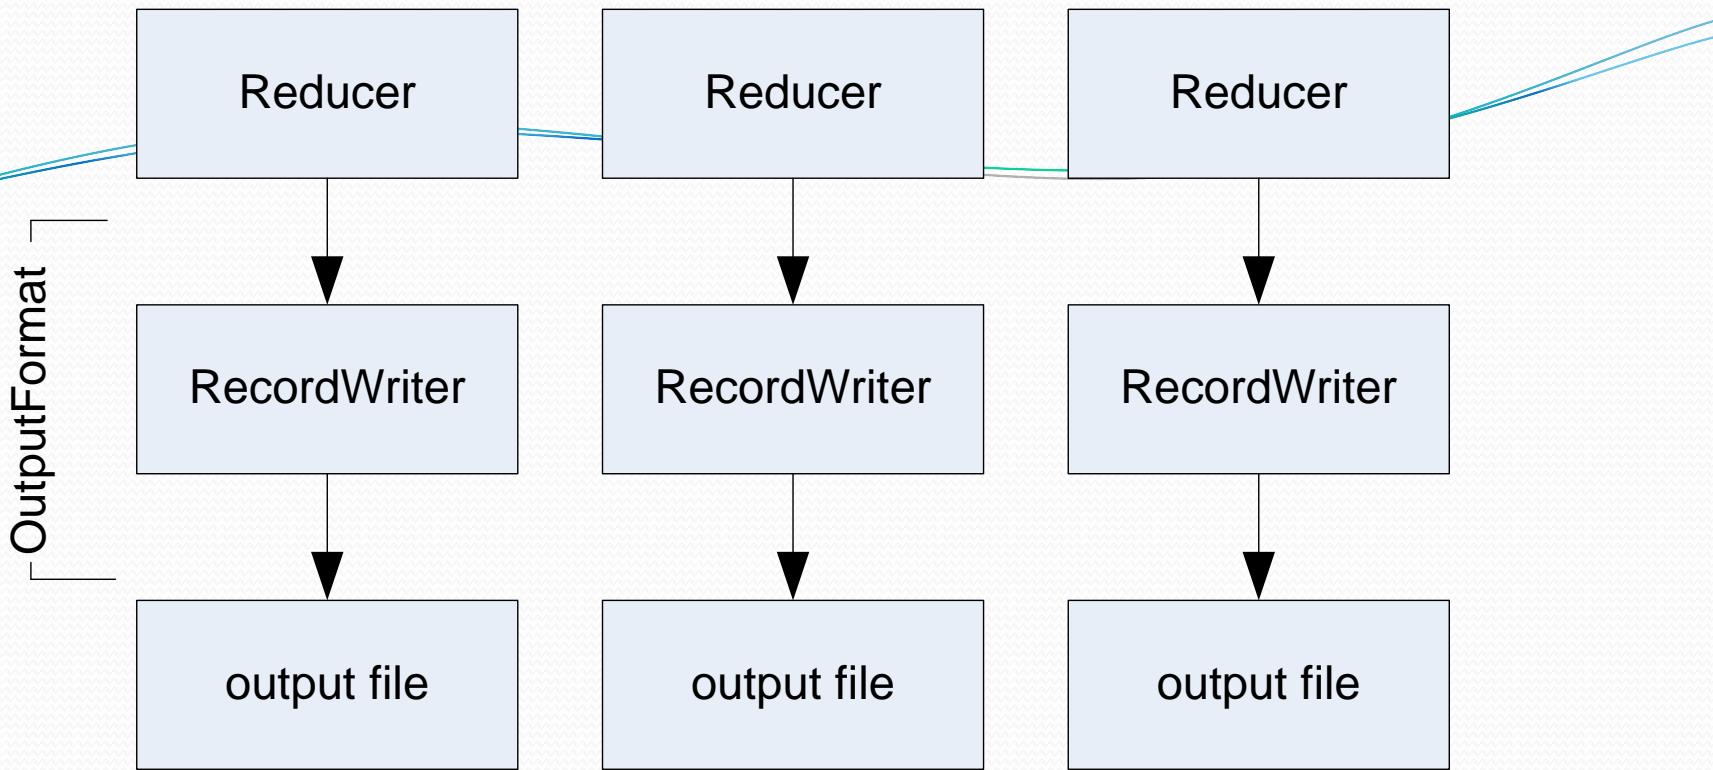
MapReduce Flow - Step1



MapReduce Flow - Step2



MapReduce flow - Step 3



Log processing

```
10.77.25.72 [16/Feb/2009:14:29:27] http://www.google.com/
```

```
10.77.25.72 [16/Feb/2009:14:29:27] http://www.rediff.com/
```

....

```
10.88.45.101 [16/Feb/2009:16:58:13] http://www.cnn.com/
```

```
10.88.45.101 [16/Feb/2009:16:58:13] http://www.yahoo.com/
```

```
10.77.45.82 [17/Feb/2009:09:35:42] http://news.cnet.com/
```

```
10.88.45.101 [18/Feb/2009:16:03:07] http://www.cnn.com/
```

....

```
10.88.45.101 [18/Feb/2009:16:03:07] http://www.yahoo.com/
```

```
10.88.88.69 [09/Jun/2009:15:43:50] http://www.cnn.com/
```

```
10.88.88.69 [11/Jun/2009:16:11:25] http://news.cnet.com/
```

```
10.88.88.69 [11/Jun/2009:16:11:25] http://www.cnn.com/
```

....

```
10.77.224.66 [12/Jun/2009:14:41:54] http://www.rediff.com/
```

```
10.77.6.54 [12/Jun/2009:15:05:31] http://www.yahoo.com/
```

```
10.88.88.69 [19/Jun/2009:16:31:11] http://www.cnn.com/
```

```
10.77.222.22 [26/Jun/2009:15:07:46] http://www.yahoo.com/
```

```
10.77.222.22 [26/Jun/2009:15:07:50] http://www.cnn.com/
```

```
10.77.222.22 [26/Jun/2009:15:11:56] http://www.yahoo.com/
```

```
10.88.88.69 [26/Jun/2009:15:48:15] http://www.cnn.com/
```

```
10.88.88.69 [02/Jul/2009:16:11:51] http://www.yahoo.com/
```

....

....

Log file may be in
Terabytes

Find Url
Frequencies

Log processing

Key: Offset into the file

Value: 10.88.45.101 [16/Feb/2009:16:58:13]

http://www.cnn.com/

Input to Mapper

Key: http://www.cnn.com/

Value: 1

Output of Mapper

Key: http://www.cnn.com/

Value: [1,1,1,1,1/.....] or [n1, n2, n3 ...] if combiner is used

Input to Reducer

Key1: http://www.cnn.com/

Value1: 1000

Key2: http://www.yahoo.com/

Value2: 5000

Output from Reducer

Search

- **Input:** (lineNumber, line) records
- **Output:** lines matching a given pattern
- **Map:**

```
if(line matches pattern):  
    output(line)
```
- **Reduce:** identify function
 - Alternative: no reducer (map-only job)

Sort

- **Input:** (key, value) records
 - **Output:** same records, sorted by key
-
- **Map:** identity function
 - **Reduce:** identify function
-
- **Trick:** Pick partitioning function h such that
 $k_1 < k_2 \Rightarrow h(k_1) < h(k_2)$

Inverted Index

- **Input:** (filename, text) records
- **Output:** list of files containing each word
- **Map:**

```
foreach word in text.split():
    output(word, filename)
```
- **Combine:** uniquify filenames for each word
- **Reduce:**

```
def reduce(word, filenames):
    output(word, sort(filenames))
```

Most Popular words

- **Input:** (filename, text) records
- **Output:** the 100 words occurring in most files
- Two-stage solution:
 - **Job 1:**
 - Create inverted index, giving (word, list(file)) records
 - **Job 2:**
 - Map each (word, list(file)) to (count, word)
 - Sort these records by count as in sort job
- Optimizations:
 - Map to (word, 1) instead of (word, file) in Job 1
 - Estimate count distribution in advance by sampling

Recommendation System

The screenshot shows the LinkedIn interface with the title "People You May Know". At the top, there are links for Home, Profile, Contacts, Groups, Jobs, Inbox (112), Companies, News, More, People, and a search bar. Below this, a banner says "See people from different parts of your professional life". A grid of cards displays recommendations:

- Osmania University**
- Mahindra Satyam**
- CGI**
- Jawaharlal Nehru Technolog...**
- HCL**
- Andhra University**
- Infosys**

chandra shekhar
Asst.Manager at Jubilant Life Sciences Ltd
India

Vikas Kumar 2nd
Testing Practice Manager at Cognizant Technology Solutions
Hartford, Connecticut Area

Madhavi Lakkaraju 2nd
Sr Project Manager at CGI Information Systems and Management Consultants
Hyderabad Area, India

Subbu Nimmagadda 3rd
Head - Sales and Recruitment at GSS Infotech
Greater New York City Area

Connect | **56 shared connections**

Connect | **42 shared connections**

Connect | **30 shared connections**

Input: 50M users and their connections

TB's of input data

Output: List(User, List(Top N probable connections))

Insight: You might know Vikas if a lot of people you know, know Vikas

Recommendation System

// Input: UserList: *List (username, List(Connections))*

For each user “U” in UserList: //50 Million users

 For each connection “X” in connections(U): // 100 Connections

 For each connection “Y” in connections(X): // 100 Connections

 if (“Y” not in Connections(U)): //**500 Billion iterations**

 Count(U, Y)++;

Sort (U, Y) order by Count (U, Y)

Store top N as (U, {Yo, Y1,, Yn})

Recommendation System – MapReduce

```
"Ankam" => ["Ravi", "Kiran", "Sameer", "Sandeep"]  
"Ravi" => ["Ankam", "Tushar", "Mukund", "Krishna"]  
"Sandeep" => ["Anil", "Ankam", "Mukund", "Sameer"]
```

Sample Input Record: “*Ankam*” => [“*Ravi*, “*Kiran*,” “*Sameer*”,]

Job1:

- Map:
 - Input: (Person, Connection_list)
 - Output: PersonPair: ((Person, Friend1), 0) and ((Friend1, Friend2), 1)
- Reduce:
 - Input: (Person, Friend1), List (degree)
 - Sum all frequencies for Person_Pair – Ignore if “0”
 - Output: ((Person1, Person2), Path_count)

Sample Output Record: (“*Ankam*”, “*Mukund*”) => 10

Recommendation System

"Ankam" => ["Ravi", "Kiran", "Sameer", "Sandeep"]

(("Ankam", "Ravi"), 0)
((("Ankam", "Kiran"), 0)
((("Ankam", "Sameer"), 0)
((("Ankam", "Sandeep"), 0)

((("Ravi", "Kiran"), 1)
((("Ravi", "Sameer"), 1)
((("Ravi", "Sandeep"), 1)

((("Kiran", "Sameer"), 1)
((("Kiran", "Sandeep"), 1)

((("Sameer", "Sandeep"), 1)

"Ravi" => ["Ankam", "Tushar", "Mukund", "Krishna"]

((("Ravi", "Ankam"), 0)
((("Ravi", "Tushar"), 0)
((("Ravi", "Mukund"), 0)
((("Ravi", "Krishna"), 0)

((("Ankam", "Tushar"), 1)
((("Ankam", "Mukund"), 1)
((("Ankam", "Krishna"), 1)

((("Tushar", "Mukund"), 1)
((("Tushar", "Krishna"), 1)

((("Mukund", "Krishna"), 1)

"Sandeep" => ["Anil", "Ankam", "Mukund", "Sameer"]

((("Sandeep", "Anil"), 0)
((("Sandeep", "Ankam"), 0)
((("Sandeep", "Mukund"), 0)
((("Sandeep", "Sameer"), 0)

((("Anil", "Ankam"), 1)
((("Anil", "Mukund"), 1)
((("Anil", "Sameer"), 1)

((("Ankam", "Mukund"), 1)
((("Ankam", "Sameer"), 1)

((("Mukund", "Sameer"), 1)

...
((("Ankam", "Sameer"), [0, 1])
((("Ankam", "Sandeep"), [0])
((("Ankam", "Mukund"), [1, 1])
((("Ankam", "Tushar"), [1]))
...

...
...
((Ankam, Kiran), [0])
...

...
((("Ankam", "Mukund"), 2)
((("Ankam", "Tushar"), 1))
...

Output of
Mappers

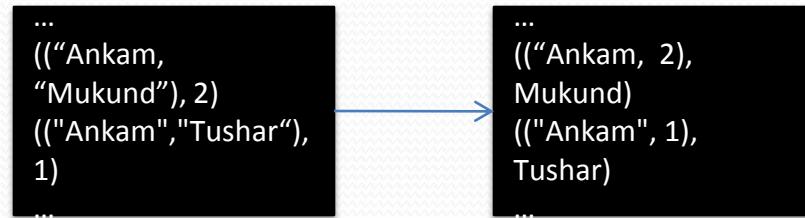
Input to
Reducers

Output of
Reducers

Recommendation System

Job2:

- Map: Just rearrange the records
 - Input: ((Person_Pair), path_count)
 - Output: ((Person1, path_count), Person2)
- Partition on first element of the key
- Reduce:
 - Input: (Connection_count_pair, candidate_list)
 - All the connections for a person are sorted by path count. Just take the top N and output
 - Output: Person, List (Friend, Count)



Since Partitioning is on first element of Key – All records for a Person will be given to one Reducer

TF-IDF

- Term frequency–Inverse document frequency
- TF-IDF measure used to evaluate how important a word is to a document in a collection
- Used in information retrieval and text mining

TF : Number of times a given terms appears in a document

$$\text{(Term frequency)} / \text{(Size of the doc)}$$

IDF: Measures the general importance of the term:

$$\text{Log} \left(\frac{\text{Total number of docs}}{\text{Number of docs in which term appears}} \right)$$

Why IDF?

- Query: Details on Hadoop: 194,000 results (194K)
 - Details: 1,300,000,000 (1.3B results)
 - On: 11,340,000,000 (11.34B results)
 - Hadoop: 1,810,000 (1.8M results)

Rare words (for a given doc collection) in query should be given more weightage

TF-IDF Algo

// Input: List(Docs)

Foreach Doc in Input :

 Tokenize and Get the words

 For each Word:

 Count number of times word occurs

// Result: TF = ((Word, DocID), TF)

Group by words and count

// Result: List (word, (List(DocID), count))

Compute IDF and then Compute TF*IDF

TF-IDF MapReduce

- **Job1: Word Frequency in Doc and Total word count for docs**

- Mapper:
 - Input: (docname, content)
 - Output: [(word, docname), (n, N)]
- Reducer:
 - Identity Reducer

n : Number of times a term appears in given Doc

N : Total number of words in given doc

Doc1
... Hadoop ...
Details.... On

Doc2
.... Details on
Cluster ...

Docn
Hadoop...
On....

Input

[(Hadoop, Doc1), (1, 1000)]
[(Details, Doc1), (10, 1000)]
[(On, Doc1), (100, 1000)]
[(Details, Doc2), (200, 2000)]
....

[(Hadoop, Doc10), (5, 5000)]
[(Details, Doc10), (50, 5000)]
[(On, Doc25), (500, 5000)]
[(Details, Doc25), (400, 4000)]
....

Output

TF-IDF MapReduce

- **Job2: Word Frequency in corpus**

- Mapper:

- Input: ((word, docname), (n,N))
 - Output: (word, (docname, n, N, 1))

- Reducer:

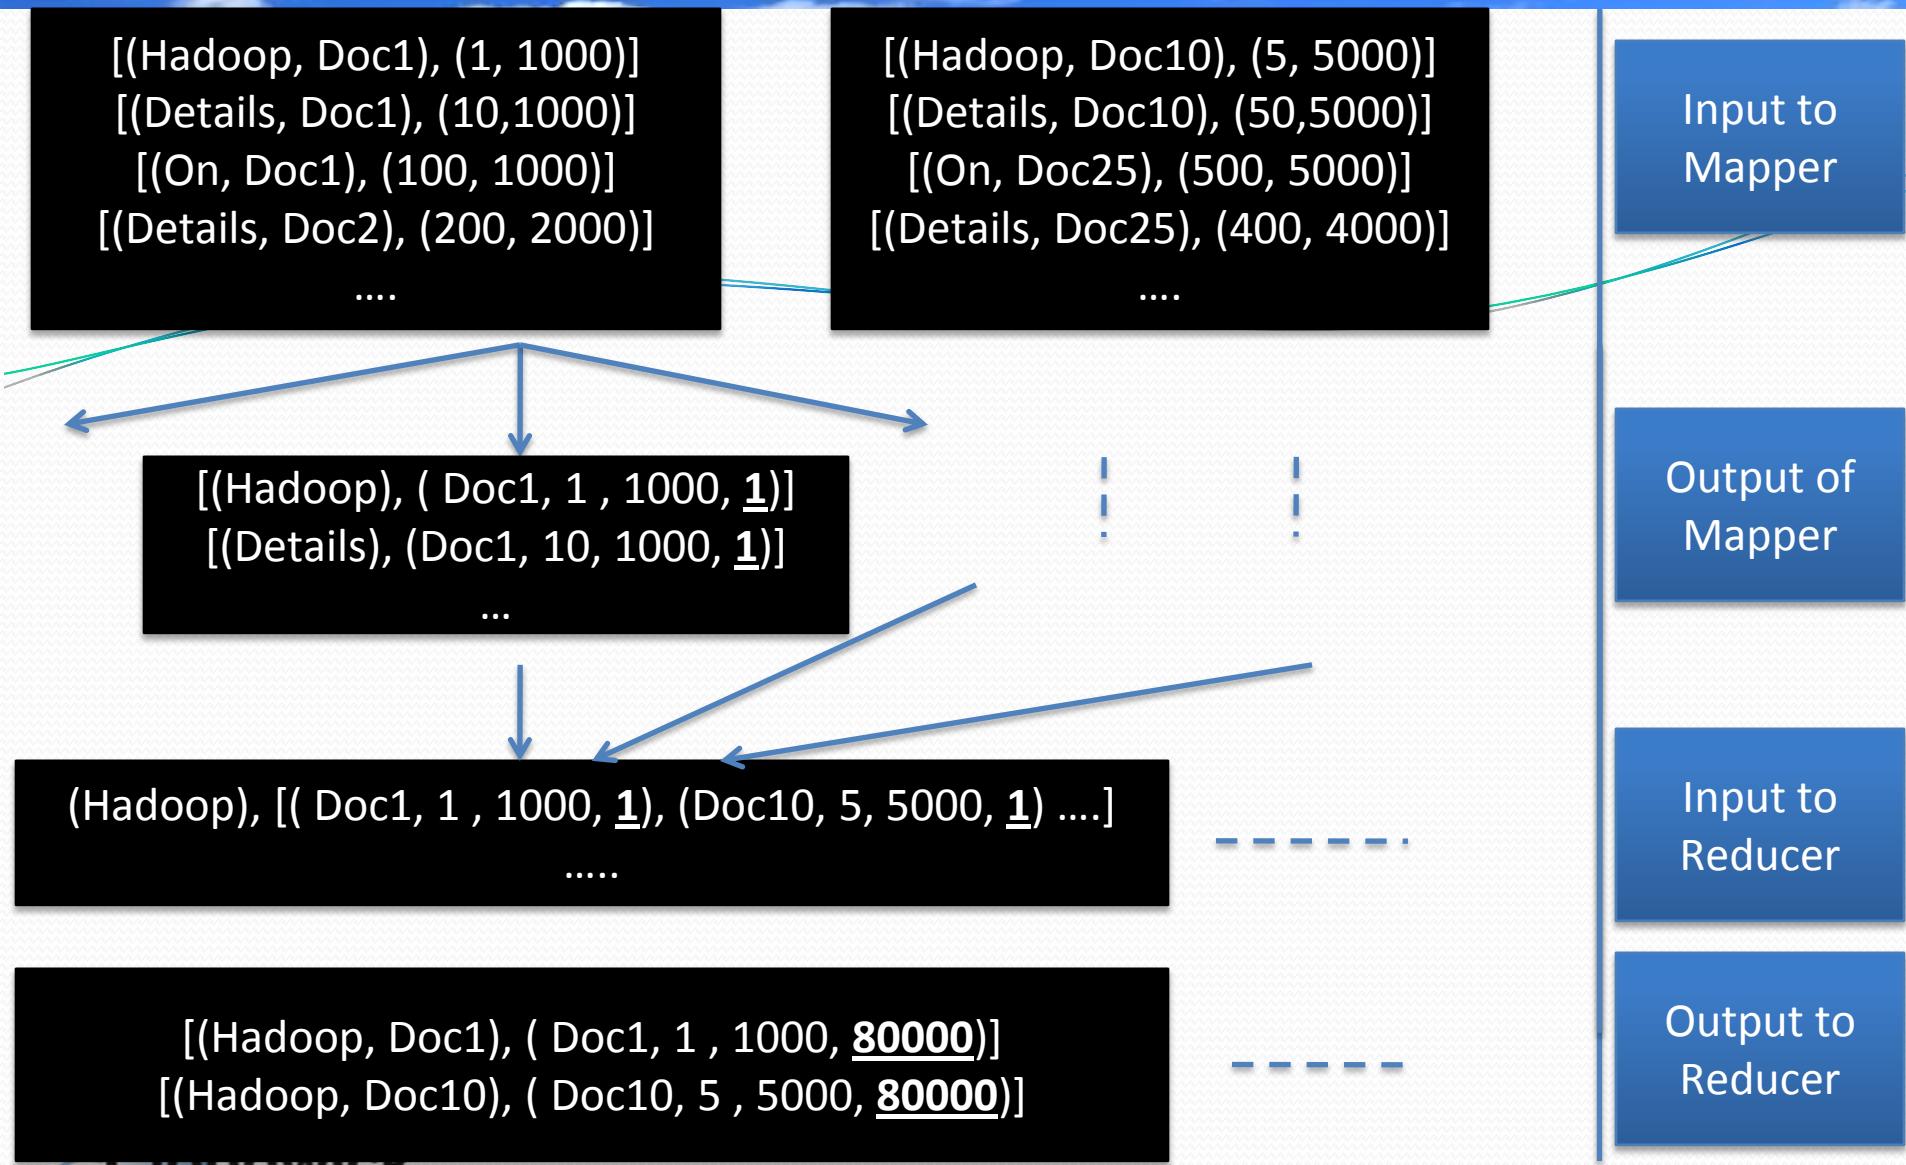
- Input: (word, [(docname₁, n₁, N₁, 1), (docname₂, n₂, N₂, 1)])
 - Add-up count for word in corpus
 - Output: (word, document), (n,N,m))

n : Number of times a term appears in given Doc

N : Total number of words in given doc

m : Number of documents in which word occurs

TF-IDF MapReduce



TF-IDF MapReduce

- **Job3: Calculate TF-IDF**

- Mapper:
 - Input: (word, document), (n,N,m)
 - Output: ((word, docname), TF*IDF)
- Reducer:
 - Do Nothing: Identity Function

TF = (Term frequency) / (Size of the doc)

IDF = Log (Total number of docs / Number of docs in which term appears)

n : Number of times a term appears in given Doc
N : Total number of words in given doc
m : Number of documents in which word occurs

Questions?

